

```
In [32]: from google.cloud import bigquery

# Use raw string for the file path
key_path = r"C:\Phoenix Portfolio\phoenix-portfolio-461608-e794abe0cf71.json"

# Initialize the BigQuery client
client = bigquery.Client.from_service_account_json(key_path)

# Query table
query = """
    SELECT *
    FROM `bigquery-public-data.ml_datasets.credit_card_default`
    LIMIT 10000
    """

# Run the query and create a DataFrame
df = client.query(query).to_dataframe()

# Display the first few rows
print(df.head())
```

	id	limit_balance	sex	education_level	marital_status	age	pay_0	\
0	27502.0	80000.0	1	6	1	54.0	0.0	
1	26879.0	200000.0	1	4	1	49.0	0.0	
2	18340.0	20000.0	2	6	2	22.0	0.0	
3	13692.0	260000.0	2	4	2	33.0	0.0	
4	20405.0	150000.0	1	4	2	32.0	0.0	

	pay_2	pay_3	pay_4	...	bill_amt_5	bill_amt_6	pay_amt_1	pay_amt_2	\
0	0.0	0.0	0.0	...	26210.0	17643.0	2545.0	2208.0	
1	0.0	0.0	0.0	...	50235.0	48984.0	1689.0	2164.0	
2	0.0	0.0	0.0	...	500.0	0.0	4641.0	1019.0	
3	0.0	0.0	0.0	...	30767.0	29890.0	5000.0	5000.0	
4	0.0	0.0	-1.0	...	143375.0	146411.0	4019.0	146896.0	

	pay_amt_3	pay_amt_4	pay_amt_5	pay_amt_6	default_payment_next_month	\
0	1336.0	2232.0	542.0	348.0		1
1	2500.0	3480.0	2500.0	3000.0		0
2	900.0	0.0	1500.0	0.0		1
3	1137.0	5000.0	1085.0	5000.0		0
4	157436.0	4600.0	4709.0	5600.0		0

	predicted_default_payment_next_month
0	[{'tables': {'score': 0.8667634129524231, 'val...
1	[{'tables': {'score': 0.9351968765258789, 'val...
2	[{'tables': {'score': 0.8572560548782349, 'val...
3	[{'tables': {'score': 0.9690881371498108, 'val...
4	[{'tables': {'score': 0.9349926710128784, 'val...

[5 rows x 26 columns]

```
In [33]: #Initial Inspection
print("Dataset shape:", df.shape)
df.head()
print("\nInfo:")
print(df.info()) #understand the purpose of each function
```

```
print("\nSummary Statistics:")
print(df.describe())
print("\nMissing Values:")
print(df.isnull().sum())
print(type(df))
for col in df.columns:
    print(col, type(df[col].iloc[0]))
df_no_ndarray = df.drop(columns=["predicted_default_payment_next_month"])
print("Duplicated Rows:", df_no_ndarray.duplicated().sum())
print("\nData Types:")
print(df.dtypes)
```

Dataset shape: (2965, 26)

Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 2965 entries, 0 to 2964

Data columns (total 26 columns):

#	Column	Non-Null Count	Dtype
0	id	2965 non-null	float64
1	limit_balance	2965 non-null	float64
2	sex	2965 non-null	object
3	education_level	2965 non-null	object
4	marital_status	2965 non-null	object
5	age	2965 non-null	float64
6	pay_0	2965 non-null	float64
7	pay_2	2965 non-null	float64
8	pay_3	2965 non-null	float64
9	pay_4	2965 non-null	float64
10	pay_5	2965 non-null	object
11	pay_6	2965 non-null	object
12	bill_amt_1	2965 non-null	float64
13	bill_amt_2	2965 non-null	float64
14	bill_amt_3	2965 non-null	float64
15	bill_amt_4	2965 non-null	float64
16	bill_amt_5	2965 non-null	float64
17	bill_amt_6	2965 non-null	float64
18	pay_amt_1	2965 non-null	float64
19	pay_amt_2	2965 non-null	float64
20	pay_amt_3	2965 non-null	float64
21	pay_amt_4	2965 non-null	float64
22	pay_amt_5	2965 non-null	float64
23	pay_amt_6	2965 non-null	float64
24	default_payment_next_month	2965 non-null	object
25	predicted_default_payment_next_month	2965 non-null	object

dtypes: float64(19), object(7)

memory usage: 602.4+ KB

None

Summary Statistics:

	id	limit_balance	age	pay_0	pay_2 \
count	2965.000000	2965.000000	2965.000000	2965.000000	2965.000000
mean	14945.556155	163369.308600	35.193255	0.005059	-0.122428
std	8700.288152	125030.415472	9.109439	1.114395	1.180784
min	29.000000	10000.000000	21.000000	-2.000000	-2.000000
25%	7499.000000	50000.000000	28.000000	-1.000000	-1.000000
50%	14782.000000	140000.000000	34.000000	0.000000	0.000000
75%	22571.000000	230000.000000	41.000000	0.000000	0.000000
max	29995.000000	800000.000000	69.000000	8.000000	7.000000

	pay_3	pay_4	bill_amt_1	bill_amt_2	bill_amt_3 \
count	2965.000000	2965.000000	2965.000000	2965.000000	2965.000000
mean	-0.141653	-0.185160	52118.305228	50649.153120	48239.757504
std	1.183630	1.178322	72328.670541	70785.001588	68145.710745
min	-2.000000	-2.000000	-11545.000000	-67526.000000	-25443.000000
25%	-1.000000	-1.000000	3958.000000	3390.000000	3302.000000
50%	0.000000	0.000000	24257.000000	23111.000000	21520.000000

75%	0.000000	0.000000	69852.000000	67827.000000	63023.000000
max	7.000000	8.000000	613860.000000	512650.000000	578971.000000

	bill_amt_4	bill_amt_5	bill_amt_6	pay_amt_1 \
count	2965.000000	2965.000000	2965.000000	2965.000000
mean	44089.683305	40956.080607	39773.072513	6348.902867
std	61907.454056	58271.904751	57303.488981	20885.735336
min	-46627.000000	-46627.000000	-73895.000000	0.000000
25%	2582.000000	1958.000000	1430.000000	1013.000000
50%	19894.000000	18814.000000	18508.000000	2234.000000
75%	58622.000000	53373.000000	52287.000000	5087.000000
max	488808.000000	441981.000000	436172.000000	493358.000000

	pay_amt_2	pay_amt_3	pay_amt_4	pay_amt_5 \
count	2.965000e+03	2965.000000	2965.000000	2965.000000
mean	6.272494e+03	5150.497133	4561.376054	4913.286678
std	2.887967e+04	14287.079982	13281.499599	16734.340778
min	0.000000e+00	0.000000	0.000000	0.000000
25%	9.900000e+02	477.000000	313.000000	323.000000
50%	2.175000e+03	1994.000000	1600.000000	1646.000000
75%	5.000000e+03	4500.000000	4000.000000	4021.000000
max	1.227082e+06	199209.000000	202076.000000	388071.000000

	pay_amt_6
count	2965.000000
mean	5382.701518
std	17275.953029
min	0.000000
25%	173.000000
50%	1615.000000
75%	4081.000000
max	403500.000000

Missing Values:

id	0
limit_balance	0
sex	0
education_level	0
marital_status	0
age	0
pay_0	0
pay_2	0
pay_3	0
pay_4	0
pay_5	0
pay_6	0
bill_amt_1	0
bill_amt_2	0
bill_amt_3	0
bill_amt_4	0
bill_amt_5	0
bill_amt_6	0
pay_amt_1	0
pay_amt_2	0
pay_amt_3	0
pay_amt_4	0

```

pay_amt_5                                0
pay_amt_6                                0
default_payment_next_month               0
predicted_default_payment_next_month     0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
id <class 'numpy.float64'>
limit_balance <class 'numpy.float64'>
sex <class 'str'>
education_level <class 'str'>
marital_status <class 'str'>
age <class 'numpy.float64'>
pay_0 <class 'numpy.float64'>
pay_2 <class 'numpy.float64'>
pay_3 <class 'numpy.float64'>
pay_4 <class 'numpy.float64'>
pay_5 <class 'str'>
pay_6 <class 'str'>
bill_amt_1 <class 'numpy.float64'>
bill_amt_2 <class 'numpy.float64'>
bill_amt_3 <class 'numpy.float64'>
bill_amt_4 <class 'numpy.float64'>
bill_amt_5 <class 'numpy.float64'>
bill_amt_6 <class 'numpy.float64'>
pay_amt_1 <class 'numpy.float64'>
pay_amt_2 <class 'numpy.float64'>
pay_amt_3 <class 'numpy.float64'>
pay_amt_4 <class 'numpy.float64'>
pay_amt_5 <class 'numpy.float64'>
pay_amt_6 <class 'numpy.float64'>
default_payment_next_month <class 'str'>
predicted_default_payment_next_month <class 'numpy.ndarray'>
Duplicated Rows: 0

```

Data Types:

id	float64
limit_balance	float64
sex	object
education_level	object
marital_status	object
age	float64
pay_0	float64
pay_2	float64
pay_3	float64
pay_4	float64
pay_5	object
pay_6	object
bill_amt_1	float64
bill_amt_2	float64
bill_amt_3	float64
bill_amt_4	float64
bill_amt_5	float64
bill_amt_6	float64
pay_amt_1	float64
pay_amt_2	float64
pay_amt_3	float64

```


pay_amt_4                                float64
pay_amt_5                                float64
pay_amt_6                                float64
default_payment_next_month                object
predicted_default_payment_next_month      object
dtype: object

```

```

In [34]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(6, 4))
sns.countplot(x='sex', hue='default_payment_next_month', data=df)
plt.title("Default Payment by Sex")
plt.xlabel("sex (1 = Male, 2 = Female)")
plt.ylabel("Count")
plt.legend(title='Default')
plt.show()


```




```

In [35]: plt.figure(figsize=(8, 4))
sns.countplot(x='education_level', hue='default_payment_next_month', data=df)
plt.title("Default Payment by Education Level")
plt.xlabel("Education (1 = Grad School, 2 = University, 3 = High School, 4+ = Other")
plt.ylabel("Count")
plt.legend(title='Default')
plt.show()

```



```
In [36]: plt.figure(figsize=(8, 4))
sns.countplot(x='marital_status', hue='default_payment_next_month', data=df)
plt.title("Default Payment by Marital Status")
plt.xlabel("Marriage (1 = Married, 2 = Single, 3 = Others)")
plt.ylabel("Count")
plt.legend(title='Default')
plt.show()
```



```
In [37]: import numpy as np
import pandas as pd

# Function to clean values
def clean_binary_column(val):
    # Handle Lists/arrays with a single item
```

```

if isinstance(val, (list, np.ndarray)) and len(val) == 1:
    val = val[0]
elif isinstance(val, dict) and 'value' in val:
    val = val['value']
# Try converting to numeric
try:
    return int(float(val))
except (ValueError, TypeError):
    return np.nan

# Apply cleaning
df['default_payment_next_month'] = df['default_payment_next_month'].apply(clean_bin)
df['default_payment_next_month'] = pd.to_numeric(df['default_payment_next_month'],

# Drop invalid values and cast to int
df = df.dropna(subset=['default_payment_next_month'])
df['default_payment_next_month'] = df['default_payment_next_month'].astype(int)

# Confirm it worked
print(df['default_payment_next_month'].unique())
print(df['default_payment_next_month'].dtype)

```

[1 0]

int64

```

In [38]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# 1. Accept only numeric values
df_numeric = df.select_dtypes(include=[np.number]).copy()

# 2. Drop ID column if present
if 'id' in df_numeric.columns:
    df_numeric.drop(columns=['id'], inplace=True)

# 3. Drop missing values
df_numeric.dropna(inplace=True)

# 4. Plot correlation heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(df_numeric.corr(), cmap='coolwarm', annot=False)
plt.title("Feature Correlation Heatmap")
plt.show()

# 5. Correlation with target variable
target_col = "default_payment_next_month"
if target_col in df_numeric.columns:
    correlation = df_numeric.corr()[target_col].sort_values(ascending=False)
    print(f"\nCorrelation with '{target_col}':")
    print(correlation)
else:
    print(f"\nColumn '{target_col}' not found in df_numeric.")

# 6. Drop highly correlated features (multicollinearity filter)


```



```
# Threshold can be adjusted (e.g., 0.8 or 0.9)
corr_matrix = df_numeric.corr().abs()
upper_tri = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))
to_drop = [column for column in upper_tri.columns if any(upper_tri[column] > 0.8)]

print(f"\nDropping highly correlated features:\n{to_drop}")
df_reduced = df_numeric.drop(columns=to_drop)

# 7. Box plots for outlier detection
print("\nBoxplots for numeric columns (checking for outliers)...")
for column in df_reduced.columns:
    plt.figure(figsize=(8, 4))
    sns.boxplot(data=df_reduced, x=column)
    plt.title(f"Boxplot of {column}")
    plt.tight_layout()
    plt.show()
```



Correlation with 'default_payment_next_month':

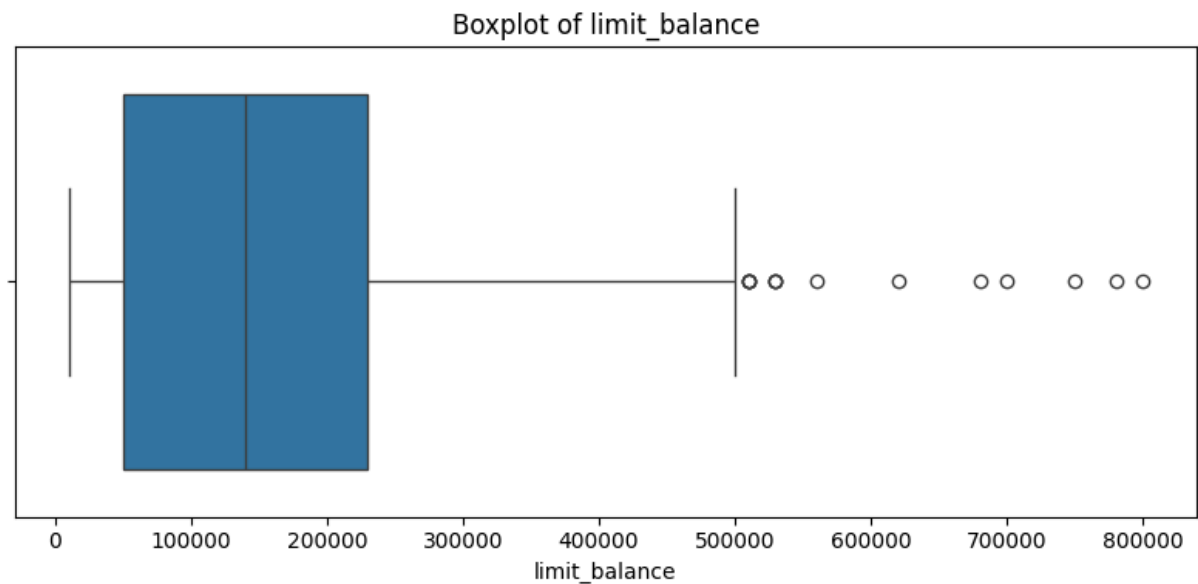
default_payment_next_month	1.000000
pay_0	0.356963
pay_2	0.288813
pay_3	0.275758
pay_4	0.255805
age	0.016905
bill_amt_6	0.006116
bill_amt_5	0.005930
bill_amt_2	0.000877
bill_amt_3	-0.001075
bill_amt_4	-0.004718
bill_amt_1	-0.005961
pay_amt_4	-0.053684
pay_amt_2	-0.059265
pay_amt_5	-0.066159
pay_amt_1	-0.066163
pay_amt_6	-0.066966
pay_amt_3	-0.090248
limit_balance	-0.161909

Name: default_payment_next_month, dtype: float64


Dropping highly correlated features:

['bill_amt_2', 'bill_amt_3', 'bill_amt_4', 'bill_amt_5', 'bill_amt_6']


Boxplots for numeric columns (checking for outliers)...




Boxplot of age




Boxplot of pay_0




Boxplot of pay_2




Boxplot of pay_3




Boxplot of pay_4




Boxplot of bill_amt_1




Boxplot of pay_amt_1




Boxplot of pay_amt_2




Boxplot of pay_amt_3




Boxplot of pay_amt_4




Boxplot of pay_amt_5



Boxplot of pay_amt_6





```
In [39]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import xgboost as xgb

# Rename target column
df.rename(columns={"default_payment_next_month": "default"}, inplace=True)

# Drop ID column
if 'id' in df.columns:
    df.drop(columns=["id"], inplace=True)

# Drop rows with missing values
df.dropna(inplace=True)

# Split features and target
X = df.drop("default", axis=1)
y = df[["default"]]

# Keep only numeric features
X = X.select_dtypes(include=["number"])

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train model
model = xgb.XGBClassifier(use_label_encoder=False, eval_metric='logloss')
model.fit(X_train_scaled, y_train)
```

```
# Predict
y_pred = model.predict(X_test_scaled)

# Evaluate
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

```
c:\Users\ADMIN\Desktop\Phoenix Analytics\venv\Lib\site-packages\xgboost\training.py:
183: UserWarning: [08:43:26] WARNING: C:\actions-runner\work\xgboost\xgboost\src\le
arner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

```
Accuracy: 0.7892074198988196
```

```
Confusion Matrix:
```

```
[[420 46]
 [ 79 48]]
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.84	0.90	0.87	466
1	0.51	0.38	0.43	127
accuracy			0.79	593
macro avg	0.68	0.64	0.65	593
weighted avg	0.77	0.79	0.78	593

```
In [40]: from sklearn.model_selection import GridSearchCV
import xgboost as xgb

# Define the model
xgb_clf = xgb.XGBClassifier(use_label_encoder=False, eval_metric='logloss')

# Define parameter grid to search
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.8, 1.0],
    'colsample_bytree': [0.8, 1.0]
}

# Set up GridSearchCV
grid_search = GridSearchCV(
    estimator=xgb_clf,
    param_grid=param_grid,
    scoring='accuracy',      # You can change to 'roc_auc', 'f1', etc.
    cv=5,                   # 5-fold cross-validation
    verbose=1,              # Shows progress
    n_jobs=-1               # Use all CPU cores
)
```



```
# Fit the grid search to the training data
grid_search.fit(X_train_scaled, y_train.values.ravel())

# Best parameters and score
print("Best Parameters:\n", grid_search.best_params_)
print("\nBest Accuracy Score:", grid_search.best_score_)
```

Fitting 5 folds for each of 108 candidates, totalling 540 fits

Best Parameters:

```
{'colsample_bytree': 1.0, 'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 100, 'subsample': 0.8}
```

Best Accuracy Score: 0.8398010215411947

```
c:\Users\ADMIN\Desktop\Phoenix Analytics\venv\Lib\site-packages\xgboost\training.py:
183: UserWarning: [08:45:26] WARNING: C:\actions-runner\work\xgboost\xgboost\src\le
arner.cc:738:
```

```
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

In [41]: `best_model = grid_search.best_estimator_`

```
# Predict
y_pred = best_model.predict(X_test_scaled)

# Evaluate
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

Accuracy: 0.8026981450252951

Confusion Matrix:

```
[[428  38]
 [ 79  48]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.92	0.88	466
1	0.56	0.38	0.45	127
accuracy			0.80	593
macro avg	0.70	0.65	0.67	593
weighted avg	0.78	0.80	0.79	593